

Nama: Rizki Aprilia Rahman

NIM: 1103213007

TUGAS BESAR UAS PEMBELAJARAN MESIN

Jobdesk:

Member 2: NLP Hugging Face Course (Chapters 1-4)

Reproduce the code from Chapters 1-4 of the NLP Hugging Face Course.

Create a detailed video tutorial explaining these chapters.


Cover key takeaways and any issues encountered.

Link Youtube: <https://youtu.be/eGtD-R5Dd8k>

```
[ ] from transformers import pipeline


classifier = pipeline("sentiment-analysis")
classifier("I've been waiting for a HuggingFace course my whole life.")

[{'label': 'POSITIVE', 'score': 0.9598047137260437}]
```

 No model was supplied, defaulted to distilbert/distilbert-base-uncased-finetuned-sst-2-english and revision 714eb0f (<https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english>)
Using a pipeline without specifying a model name and revision in production is not recommended.
Device set to use cpu
[{'label': 'POSITIVE', 'score': 0.9598047137260437}]

kode ini mendemonstrasikan bagaimana menggunakan model analisis sentimen dari Hugging Face untuk mengevaluasi teks secara otomatis.

```
[ ] classifier(
    ["I've been waiting for a HuggingFace course my whole life.", "I hate this so much!"]
)
[{'label': 'POSITIVE', 'score': 0.9598047137260437},
 {'label': 'NEGATIVE', 'score': 0.9994558095932007}]
```


 [{'label': 'POSITIVE', 'score': 0.9598047137260437},
{'label': 'NEGATIVE', 'score': 0.9994558095932007}]


kode ini menunjukkan cara menganalisis beberapa kalimat sekaligus menggunakan pipeline sentimen dari Hugging Face, dan model dapat dengan akurat menentukan sentimen dari masing-masing kalimat.


```
[ ] # Zero-shot classification
from transformers import pipeline


classifier = pipeline("zero-shot-classification")
classifier(
    "This is a course about the Transformers library",
    candidate_labels=["education", "politics", "business"],
)


{'sequence': 'This is a course about the Transformers library',
 'labels': ['education', 'business', 'politics'],
 'scores': [0.8445963859558105, 0.111976258456707, 0.043427448719739914]}
```


 No model was supplied, defaulted to facebook/bart-large-mnli and revision d7645e1 (<https://huggingface.co/facebook/bart-large-mnli>). Using a pipeline without specifying a model name and revision in production is not recommended.


config.json: 100%  1.15k/1.15k [00:00<00:00, 19.6kB/s]

model.safetensors: 100%  1.63G/1.63G [00:16<00:00, 177MB/s]

tokenizer_config.json: 100%  26.0/26.0 [00:00<00:00, 1.68kB/s]

vocab.json: 100%  899k/899k [00:00<00:00, 13.6MB/s]

merges.txt: 100%  456k/456k [00:00<00:00, 13.6MB/s]

tokenizer.json: 100%  1.36M/1.36M [00:00<00:00, 9.99MB/s]

Device set to use cpu

```
{'sequence': 'This is a course about the Transformers library',
 'labels': ['education', 'business', 'politics'],
 'scores': [0.8445963859558105, 0.111976258456707, 0.043427448719739914]}
```

kode ini menunjukkan cara menggunakan zero-shot classification untuk mengklasifikasikan teks ke dalam label yang sudah ditentukan sebelumnya, tanpa memerlukan pelatihan ulang pada model dengan label-label tersebut.

```
→ No model was supplied, defaulted to openai-community/gpt2 and revision 607a30d (https://huggingface.co/open)  
Using a pipeline without specifying a model name and revision in production is not recommended.
```

File	Progress	Size	Speed
config.json:	100%	665/665	[00:00<00:00, 9.98kB/s]
model.safetensors:	100%	548M/548M	[00:05<00:00, 103MB/s]
generation_config.json:	100%	124/124	[00:00<00:00, 2.00kB/s]
tokenizer_config.json:	100%	26.0/26.0	[00:00<00:00, 461B/s]
vocab.json:	100%	1.04M/1.04M	[00:00<00:00, 15.6MB/s]
merges.txt:	100%	456k/456k	[00:00<00:00, 3.64MB/s]
tokenizer.json:	100%	1.36M/1.36M	[00:00<00:00, 17.5MB/s]

```
Device set to use cpu  
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.  
[{'generated_text': 'In this course, we will teach you how to understand and use data flow and data  
interchange when handling user data. We will be working with one or more of the most commonly used data  
flows - data flows of various types, as seen by the HTTP'}]]
```

```
[ ] # Using any model from the Hub in a pipeline
    from transformers import pipeline

    generator = pipeline("text-generation", model="distilgpt2")
    generator(
        "In this course, we will teach you how to",
        max_length=30,
        num_return_sequences=2,
    )

    [{'generated_text': 'In this course, we will teach you how to manipulate the world and '
                        'move your mental and physical capabilities to your advantage.'},
     {'generated_text': 'In this course, we will teach you how to become an expert and '
                        'practice realtime, and with a hands on experience on both real '
                        'time and real'}]
```

```

Device set to use cpu
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncatio
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
[{'generated_text': 'In this course, we will teach you how to manipulate the world and move your mental and
physical capabilities to your advantage.'},
 {'generated_text': 'In this course, we will teach you how to become an expert and practice realtime, and
with a hands on experience on both real time and real'}}]


```

kode ini menunjukkan bagaimana model **distilgpt2** dari Hugging Face digunakan dalam pipeline untuk menghasilkan teks lanjutan dari prompt yang diberikan. Model ini menghasilkan dua teks yang berbeda setiap kali dijalankan, dengan batasan panjang teks yang dihasilkan sesuai parameter `max_length`.

```
[ ] # Mask filling
from transformers import pipeline

unmasker = pipeline("fill-mask")
unmasker("This course will teach you all about <mask> models.", top_k=2)

[{'sequence': 'This course will teach you all about mathematical models.',
  'score': 0.19619831442832947,
  'token': 30412,
  'token_str': ' mathematical'},
 {'sequence': 'This course will teach you all about computational models.',
  'score': 0.04052725434303284,
  'token': 38163,
  'token_str': ' computational'}]
```

 No model was supplied, defaulted to distilbert/distilroberta-base and revision fb53ab8 (<https://huggingface.co/distilbert/distilroberta-base>). Using a pipeline without specifying a model name and revision in production is not recommended. Some weights of the model checkpoint at distilbert/distilroberta-base were not used when initializing RobertaForMaskedLM. - This IS expected if you are initializing RobertaForMaskedLM from the checkpoint of a model trained on another task. - This IS NOT expected if you are initializing RobertaForMaskedLM from the checkpoint of a model that you expect to be initialized with this task. Device set to use cpu

```
[{'sequence': 'This course will teach you all about mathematical models.',
  'score': 0.19619831442832947,
  'token': 30412,
  'token_str': ' mathematical'},
 {'sequence': 'This course will teach you all about computational models.',
  'score': 0.04052725434303284,
  'token': 38163,
  'token_str': ' computational'}]
```

kode ini menunjukkan bagaimana menggunakan pipeline **fill-mask** untuk mengisi placeholder dalam kalimat dengan kata yang sesuai, berdasarkan konteks yang diberikan. Model menghasilkan dua kemungkinan penggantian untuk `<mask>` dengan skor probabilitas yang menunjukkan seberapa cocok kata tersebut dalam kalimat.

```

# Named entity recognition
from transformers import pipeline

ner = pipeline("ner", grouped_entities=True)
ner("My name is Sylvain and I work at Hugging Face in Brooklyn.")

[{'entity_group': 'PER', 'score': 0.99816, 'word': 'Sylvain', 'start': 11, 'end': 18},
 {'entity_group': 'ORG', 'score': 0.97960, 'word': 'Hugging Face', 'start': 33, 'end': 45},
 {'entity_group': 'LOC', 'score': 0.99321, 'word': 'Brooklyn', 'start': 49, 'end': 57}]

```

```

No model was supplied, defaulted to dbmdz/bert-large-cased-finetuned-conll03-english and revision 4c53496 (https://huggingface.co/dbmdz/bert-large-cased-finetuned-conll03-english)
Using a pipeline without specifying a model name and revision in production is not recommended.

model.safetensors: 100% 1.33G/1.33G [00:21<00:00, 27.4MB/s]

Some weights of the model checkpoint at dbmdz/bert-large-cased-finetuned-conll03-english were not used when initializing BertForTokenClassification:
- This IS expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task (with different vocabularies and token embeddings): in this case transfer learning is not enabled into the new model.
- This IS NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be initialized from without any initialization: in this case the model was not initialized properly.

tokenizer_config.json: 100% 60.0/60.0 [00:00<00:00, 3.40kB/s]

vocab.txt: 100% 213k/213k [00:00<00:00, 3.25MB/s]

Device set to use cpu
/usr/local/lib/python3.10/dist-packages/transformers/pipelines/token_classification.py:170: UserWarning: `grouped_entities` is deprecated and will be removed in a future version, use `grouped_entities=True` instead.
warnings.warn(
[{'entity_group': 'PER',
 'score': 0.99816,
 'word': 'Sylvain',
 'start': 11,
 'end': 18},
 {'entity_group': 'ORG',
 'score': 0.9796,
 'word': 'Hugging Face',
 'start': 33,
 'end': 45},
 {'entity_group': 'LOC',
 'score': 0.99321,
 'word': 'Brooklyn',
 'start': 49,
 'end': 57}]

```

kode ini menunjukkan bagaimana menggunakan pipeline **NER** untuk mendeteksi dan mengklasifikasikan entitas yang ada dalam sebuah kalimat, seperti nama orang, organisasi, dan lokasi, beserta tingkat kepercayaan model terhadap hasil tersebut.

```


# Question answering
from transformers import pipeline


question_answerer = pipeline("question-answering")
question_answerer(
    question="Where do I work?",
    context="My name is Sylvain and I work at Hugging Face in Brooklyn",
)


{'score': 0.6385916471481323, 'start': 33, 'end': 45, 'answer': 'Hugging Face'}


```


No model was supplied, defaulted to distilbert/distilbert-base-cased-distilled-squad and revision
 Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100%  473/473 [00:00<00:00, 7.66kB/s]

model.safetensors: 100%  261M/261M [00:03<00:00, 68.8MB/s]

tokenizer_config.json: 100%  49.0/49.0 [00:00<00:00, 919B/s]

vocab.txt: 100%  213k/213k [00:00<00:00, 3.24MB/s]

tokenizer.json: 100%  436k/436k [00:00<00:00, 5.62MB/s]

Device set to use cpu

{'score': 0.6385916471481323, 'start': 33, 'end': 45, 'answer': 'Hugging Face'}


Secara keseluruhan, kode ini menunjukkan bagaimana menggunakan pipeline **question answering** untuk menjawab pertanyaan berdasarkan konteks yang diberikan, serta bagaimana model memberikan skor kepercayaan terhadap jawabannya.


```


[{'summary_text': ' America has changed dramatically during recent years . The '
                  'number of engineering graduates in the U.S. has declined in '
                  'traditional engineering disciplines such as mechanical, civil '
                  ', electrical, chemical, and aeronautical engineering . Rapidly '
                  'developing economies such as China and India, as well as other '
                  'industrial countries in Europe and Asia, continue to encourage '
                  'and advance engineering .'}]


```


No model was supplied, defaulted to sshleifer/distilbart-cnn-12-6 and revision a4f8f3e (<https://huggingface.co/sshleifer/distilbart-cnn-12-6>) and revision a4f8f3e (<https://huggingface.co/sshleifer/distilbart-cnn-12-6>)
 Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100%  1.80k/1.80k [00:00<00:00, 32.4kB/s]

pytorch_model.bin: 100%  1.22G/1.22G [00:13<00:00, 173MB/s]

tokenizer_config.json: 100%  26.0/26.0 [00:00<00:00, 566B/s]

vocab.json: 100%  899k/899k [00:00<00:00, 12.8MB/s]

merges.txt: 100%  456k/456k [00:00<00:00, 5.81MB/s]

Device set to use cpu

[{'summary_text': ' America has changed dramatically during recent years . The number of engineering graduates in the U.S. has declined in traditional engineering disciplines such as mechanical, civil , electrical, chemical, and aeronautical engineering . Rapidly developing economies such as China and India, as well as other industrial countries in Europe and Asia, continue to encourage and advance engineering .'}]

kode ini menunjukkan bagaimana menggunakan pipeline **summarization** untuk meringkas teks panjang menjadi versi yang lebih singkat namun tetap menyampaikan informasi penting.

```
[ ] # Translation
    from transformers import pipeline

    translator = pipeline("translation", model="Helsinki-NLP/opus-mt-fr-en")
    translator("Ce cours est produit par Hugging Face.")

    [{'translation_text': 'This course is produced by Hugging Face.'}]
```

⇒ Device set to use cpu
[{'translation_text': 'This course is produced by Hugging Face.'}]

kode ini menunjukkan bagaimana menggunakan pipeline **translation** untuk menerjemahkan teks dari satu bahasa ke bahasa lain menggunakan model dari Hugging Face.

```
[ ] # Bias and limitations
    from transformers import pipeline

    unmasker = pipeline("fill-mask", model="bert-base-uncased")
    result = unmasker("This man works as a [MASK].")
    print([r["token_str"] for r in result])

    result = unmasker("This woman works as a [MASK].")
    print([r["token_str"] for r in result])

    ['lawyer', 'carpenter', 'doctor', 'waiter', 'mechanic']
    ['nurse', 'waitress', 'teacher', 'maid', 'prostitute']
```

⇒ config.json: 100% ██████████ 570/570 [00:00<00:00, 13.3kB/s]

model.safetensors: 100% ██████████ 440M/440M [00:12<00:00, 21.0MB/s]

BertForMaskedLM has generative capabilities, as `prepare_inputs_for_generation` is explicit
 - If you're using `trust_remote_code=True`, you can get rid of this warning by loading the model with `from_pretrained(..., trust_remote_code=True, ...)`
 - If you are the owner of the model architecture code, please modify your model class so that it inherits from `BertLMHeadModel` instead of `BertForMaskedLM`
 - If you are not the owner of the model architecture class, please contact the model code owner to make this change
 Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForMaskedLM
 - This IS expected if you are initializing BertForMaskedLM from the checkpoint of a model initialized with specific instructions (e.g. with `--bf16`)
 - This IS NOT expected if you are initializing BertForMaskedLM from the checkpoint of a model that has generative capabilities (e.g. with `--do_sample`)

tokenizer_config.json: 100% ██████████ 48.0/48.0 [00:00<00:00, 955B/s]

vocab.txt: 100% ██████████ 232k/232k [00:00<00:00, 8.85MB/s]

tokenizer.json: 100% ██████████ 466k/466k [00:00<00:00, 7.70MB/s]

Device set to use cpu
 ['carpenter', 'lawyer', 'farmer', 'businessman', 'doctor']
 ['nurse', 'maid', 'teacher', 'waitress', 'prostitute']
 ['nurse', 'waitress', 'teacher', 'maid', 'prostitute']

bias dan **keterbatasan** dari model **BERT** ditunjukkan dengan membandingkan hasil pengisian kata untuk kalimat yang melibatkan gender. Hasil ini menunjukkan bahwa meskipun model seperti

BERT dapat menghasilkan teks dengan kecerdasan yang baik, model tersebut bisa saja mengandung **bias sosial** yang mencerminkan prasangka atau asumsi yang ada dalam data pelatihan. Hal ini penting untuk diperhatikan saat menggunakan model seperti ini dalam aplikasi dunia nyata.

```
# Behind the pipeline
from transformers import pipeline

classifier = pipeline("sentiment-analysis")
classifier([
    "I've been waiting for a HuggingFace course my whole life.",
    "I hate this so much!",
])
```

No model was supplied, defaulted to distilbert/distilbert-base-uncased-finetuned-sst-2-english. Using a pipeline without specifying a model name and revision in production is discouraged. Device set to use cpu




```
[{'label': 'POSITIVE', 'score': 0.9598049521446228},
 {'label': 'NEGATIVE', 'score': 0.9994558691978455}]
```

Pipeline sentiment-analysis di Hugging Face digunakan untuk secara otomatis menentukan sentimen dari teks yang diberikan, dengan memberikan label (POSITIVE atau NEGATIVE) dan skor kepercayaan model terhadap analisis tersebut.

Skor yang lebih tinggi menunjukkan tingkat kepastian model dalam menilai sentimen teks.

```
[ ] # Preprocessing with a tokenizer
from transformers import AutoTokenizer

checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
```

tokenizer_config.json: 100%  48.0/48.0 [00:00<00:00, 1.03kB/s]
config.json: 100%  629/629 [00:00<00:00, 14.0kB/s]
vocab.txt: 100%  232k/232k [00:00<00:00, 2.14MB/s]

Pada kode ini, menggunakan **tokenizer** dari Hugging Face untuk melakukan pra-pemrosesan teks sebelum digunakan dalam model.


```
[ ] ## Define the raw input text

raw_inputs = [
    "I've been waiting for a HuggingFace course my whole life.",
    "I hate this so much!",
]
inputs = tokenizer(raw_inputs, padding=True, truncation=True, return_tensors="pt")
print(inputs)
```

```
{'input_ids': tensor([[ 101, 1045, 1005, 2310, 2042, 3403, 2005, 1037, 17662, 12172,
 2607, 2026, 2878, 2166, 1012, 102],
 [ 101, 1045, 5223, 2023, 2061, 2172, 999, 102, 0, 0,
 0, 0, 0, 0, 0, 0]]), 'attention_mask': tensor([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
 1, 1, 1, 1, 1, 1],
 [1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0]])}
```

Pada kode ini, melakukan pra-pemrosesan teks menggunakan tokenizer yang telah dimuat sebelumnya.

```
[ ] pip install torch
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.5.1+cu121)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch) (3.16.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch) (3.4.2)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch) (3.1.4)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch) (2024.10.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch) (2.1.5)
```

Install torch agar bisa running setelahnya.

```
[ ] import torch
from transformers import AutoTokenizer, AutoModel

# The error NameError: name 'tensor' is not defined occurs because the tensor object is not imported or def.
# Key Fixes:
# Use torch.tensor explicitly: The tensor function must be accessed as torch.tensor because it belongs to tl
# Ensure AutoTokenizer is loaded: The tokenizer must be loaded to preprocess raw_inputs.
# Output Explanation:
# inputs will display the input_ids and attention_mask tensors created by the tokenizer.
# The torch.tensor section mimics what the tokenizer outputs to demonstrate how tensors are structured.
```

```
{'input_ids': tensor([[ 101, 1045, 1005, 2310, 2042, 3403, 2005, 1037, 17662, 12172,
 2607, 2026, 2878, 2166, 1012, 102],
 [ 101, 1045, 5223, 2023, 2061, 2172, 999, 102, 0, 0,
 0, 0, 0, 0, 0, 0]]), 'attention_mask': tensor([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
 1, 1, 1, 1, 1, 1],
 [1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0]])}
{'input_ids': tensor([[ 101, 1045, 1005, 2310, 2042, 3403, 2005, 1037, 17662, 12172,
 2607, 2026, 2878, 2166, 1012, 102],
 [ 101, 1045, 5223, 2023, 2061, 2172, 999, 102, 0, 0,
 0, 0, 0, 0, 0, 0]]), 'attention_mask': tensor([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
 1, 1, 1, 1, 1, 1],
 [1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0]])}
```

Kode ini menunjukkan cara memuat model dan tokenizer serta memproses teks melalui model tersebut.

```
[ ] # Going through the model

from transformers import AutoModel

checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
model = AutoModel.from_pretrained(checkpoint)

outputs = model(**inputs)
print(outputs.last_hidden_state.shape)

⇒ torch.Size([2, 16, 768])
```

Dalam kode ini, memuat model pre-trained dan kemudian menjalankan model tersebut pada input untuk mendapatkan output.

```
[ ] # Model heads: Making sense out of numbers

from transformers import AutoModelForSequenceClassification

checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
model = AutoModelForSequenceClassification.from_pretrained(checkpoint)
outputs = model(**inputs)

print(outputs.logits.shape)

⇒ torch.Size([2, 2])
```

Dalam kode ini, memuat model yang telah disesuaikan untuk tugas **klasifikasi urutan** (sequence classification) dan menjalankan model tersebut pada input.

```
[ ] # Postprocessing the output
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification

# Menampilkan logits output
print("Logits:")
print(outputs.logits)

# Konversi logits ke softmax probabilities (optional)
probabilities = torch.nn.functional.softmax(outputs.logits, dim=-1)
print("\nProbabilities:")
print(probabilities)
```

```
⇒ Logits:
tensor([[ -1.5607,  1.6123],
        [ 4.1692, -3.3464]], grad_fn=<AddmmBackward0>)

Probabilities:
tensor([[4.0195e-02, 9.5980e-01],
        [9.9946e-01, 5.4418e-04]], grad_fn=<SoftmaxBackward0>)
```

Dalam kode ini, melakukan **postprocessing** pada output dari model untuk mengubah logits menjadi probabilitas menggunakan fungsi **softmax**.

Kendala: Masalah muncul ketika mencoba menampilkan atau mencetak hasil tensor yang berisi operasi yang memiliki `grad_fn=<AddmmBackward>`, yang mengindikasikan bahwa tensor tersebut terkait dengan operasi dalam grafik komputasi PyTorch. Ini sering terjadi ketika mencoba langsung mencetak output tensor tanpa mengkonversinya terlebih dahulu menjadi format yang dapat dibaca.

Solusi:

1. **Menampilkan Logits:** Gunakan `print(outputs.logits)` untuk menampilkan tensor logits secara langsung. Jika Anda mendapatkan pesan kesalahan terkait `grad_fn`, ini berarti tensor masih memiliki referensi ke operasi di dalam grafik komputasi, dan Anda mungkin perlu menghilangkan grafik komputasi tersebut.
2. **Softmax:** Untuk mengubah logits menjadi probabilitas, gunakan `torch.nn.functional.softmax(outputs.logits, dim=-1)` untuk menghindari kesalahan format dan mengkonversi tensor logits ke probabilitas yang bisa dipahami.

```
[ ] import torch

    predictions = torch.nn.functional.softmax(outputs.logits, dim=-1)
    print(predictions)

⇒ tensor([[4.0195e-02, 9.5980e-01],
          [9.9946e-01, 5.4418e-04]], grad_fn=<SoftmaxBackward0>)
```

Kode

Kode ini akan mengubah tensor logits (hasil dari model) menjadi probabilitas dengan menggunakan fungsi softmax. Fungsi softmax akan mengubah nilai logits menjadi distribusi probabilitas dengan memastikan bahwa jumlah seluruh probabilitas adalah 1.

```
[ ] model.config.id2label

⇒ {0: 'NEGATIVE', 1: 'POSITIVE'}
```

```
[ ] # Creating a Transformer

from transformers import BertConfig, BertModel

# Building the config
config = BertConfig()

# Building the model from the config
model = BertModel(config)
print(config)
```

```
⇒ BertConfig {
  "_attn_implementation_autoset": true,
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "transformers_version": "4.47.1",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 30522
}
```

Kode ini membuat dan mencetak konfigurasi serta model BERT menggunakan pustaka transformers dari Hugging Face.

```
[ ] # Different loading methods
```

```
from transformers import BertConfig, BertModel
```

```
config = BertConfig()
```

```
model = BertModel(config)
```

```
# Model is randomly initialized!
```

```
[ ] from transformers import BertModel
```

```
model = BertModel.from_pretrained("bert-base-cased")
```



config.json: 100%  570/570 [00:00<00:00, 22.3kB/s]

model.safetensors: 100%  436M/436M [00:03<00:00, 132MB/s]

```
[ ] # Menyimpan model ke lokasi tertentu
```

```
model.save_pretrained("C:/Users/Asus/Downloads/Semester 7/ML")
```

```
# Menampilkan lokasi penyimpanan
```

```
print("Model telah disimpan di: C:/Users/Asus/Downloads/Semester 7/ML")
```



Model telah disimpan di: C:/Users/Asus/Downloads/Semester 7/ML

```
[ ] # Using a Transformer model for inference
```

```
sequences = ["Hello!", "Cool.", "Nice!"]
```



```
import torch
```

```
model_inputs = torch.tensor(encoded_sequences)
```

```
[ ] output = model(model_inputs)
```

Outputnya adalah daftar probabilitas untuk setiap label berdasarkan input yang diberikan.





```
[ ] #Word-based
    tokenized_text = "Jim Henson was a puppeteer".split()
    print(tokenized_text)
```

→ ['Jim', 'Henson', 'was', 'a', 'puppeteer']

Kode berikut memanfaatkan metode .split() untuk memecah teks menjadi token berbasis kata:

```
[ ] # loading and saving
    from transformers import BertTokenizer

    tokenizer = BertTokenizer.from_pretrained("bert-base-cased")
```

→ /usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>)
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
tokenizer_config.json: 100%  49.0/49.0 [00:00<00:00, 1.34kB/s]
vocab.txt: 100%  213k/213k [00:00<00:00, 2.44MB/s]
tokenizer.json: 100%  436k/436k [00:00<00:00, 4.34MB/s]
config.json: 100%  570/570 [00:00<00:00, 11.9kB/s]

Kode berikut menjelaskan cara **memuat tokenizer** yang sudah dilatih sebelumnya dari model **BERT** dan bagaimana tokenizer ini dapat digunakan untuk memproses teks.

```
[ ] from transformers import AutoTokenizer

    tokenizer = AutoTokenizer.from_pretrained("bert-base-cased")
```

[] tokenizer("Using a Transformer network is simple")

→ {'input_ids': [101, 7993, 170, 13809, 23763, 2443, 1110, 3014, 102], 'token_type_ids': [0, 0, 0, 0, 0, 0, 0, 0], 'attention_mask': [1, 1, 1, 1, 1, 1, 1, 1]}

Kode berikut menjelaskan cara **memuat tokenizer otomatis** dari model yang sudah dilatih sebelumnya menggunakan library Transformers.

```
[ ] tokenizer.save_pretrained("C:/Users/Asus/Downloads/Semester 7/ML")

➦ ('C:/Users/Asus/Downloads/Semester 7/ML/tokenizer_config.json',
   'C:/Users/Asus/Downloads/Semester 7/ML/special_tokens_map.json',
   'C:/Users/Asus/Downloads/Semester 7/ML/vocab.txt',
   'C:/Users/Asus/Downloads/Semester 7/ML/added_tokens.json',
   'C:/Users/Asus/Downloads/Semester 7/ML/tokenizer.json')
```

Perintah `tokenizer.save_pretrained("C:/Users/Asus/Downloads/Semester 7/ML")` digunakan untuk menyimpan tokenizer yang telah dimuat ke dalam folder yang ditentukan, sehingga dapat digunakan kembali tanpa harus mengunduh ulang.

```
[ ] # Encoding

from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-cased")

sequence = "Using a Transformer network is simple"
tokens = tokenizer.tokenize(sequence)

print(tokens)

➦ ['Using', 'a', 'Trans', '##former', 'network', 'is', 'simple']
```

Kode ini menggunakan tokenizer dari model "bert-base-cased" untuk memecah kalimat "Using a Transformer network is simple" menjadi token-token individu yang dipahami oleh model. Hasil dari `tokenizer.tokenize(sequence)` adalah daftar token yang sesuai dengan input tersebut, yang dalam hal ini adalah pemisahan kata dan sub-kata yang lebih kecil.

```
[ ] # From tokens to input IDs

ids = tokenizer.convert_tokens_to_ids(tokens)

print(ids)

➦ [7993, 170, 13809, 23763, 2443, 1110, 3014]
```

Kode ini mengonversi daftar token hasil dari proses tokenisasi (`tokens`) menjadi daftar ID numerik menggunakan metode `convert_tokens_to_ids`. ID ini adalah representasi numerik dari token yang sesuai dengan kosakata model, yang siap digunakan sebagai input untuk model.


```
[ ] # Decoding
    decoded_string = tokenizer.decode([7993, 170, 11303, 1200, 2443, 1110, 3014])
    print(decoded_string)
```

➡ Using a transformer network is simple

Kode ini mendekode daftar ID numerik [7993, 170, 11303, 1200, 2443, 1110, 3014] menjadi string asli menggunakan metode decode dari tokenizer. Hasilnya adalah teks yang direkonstruksi berdasarkan kosakata model.

```
[ ] # Handling multiple sequences

import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification

checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
model = AutoModelForSequenceClassification.from_pretrained(checkpoint)

sequence = "I've been waiting for a HuggingFace course my whole life."

tokens = tokenizer.tokenize(sequence)
ids = tokenizer.convert_tokens_to_ids(tokens)
input_ids = torch.tensor(ids)
# This line will fail.
model(input_ids)
```

➡

tokenizer_config.json: 100%	<div style="width: 100%; height: 10px; background-color: #28a745;"></div>	48.0/48.0 [00:00<00:00, 1.43kB/s]
config.json: 100%	<div style="width: 100%; height: 10px; background-color: #28a745;"></div>	629/629 [00:00<00:00, 14.8kB/s]
vocab.txt: 100%	<div style="width: 100%; height: 10px; background-color: #28a745;"></div>	232k/232k [00:00<00:00, 2.95MB/s]
model.safetensors: 100%	<div style="width: 100%; height: 10px; background-color: #28a745;"></div>	268M/268M [00:01<00:00, 173MB/s]

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-8-f7b4ef20cab7> in <cell line: 16>()
     14 input_ids = torch.tensor(ids)
     15 # This line will fail.
--> 16 model(input_ids)
```

Model gagal memproses input karena input yang diberikan adalah **sebuah urutan tunggal** (single sequence), sedangkan model dari 😊 Transformers secara default mengharapkan **beberapa urutan** (batch of sentences). Tokenizer secara otomatis menambahkan dimensi tambahan untuk memenuhi format input yang diharapkan oleh model, tetapi ketika kita mencoba mengimitasi proses ini secara manual, kita lupa menambahkan dimensi tambahan.

```
[ ] tokenized_inputs = tokenizer(sequence, return_tensors="pt")
    print(tokenized_inputs["input_ids"])
```

```
⇒ tensor([[ 101,  1045,  1005,  2310,  2042,  3403,  2005,  1037, 17662, 12172,
            2607,  2026,  2878,  2166,  1012,   102]])
```

Kode berikut ini menggunakan tokenizer untuk secara otomatis memproses teks dan menghasilkan ID input dalam bentuk tensor PyTorch (return_tensors="pt"), termasuk menambahkan dimensi batch yang diperlukan.

```
[ ] # Let's try again and add a new dimension:
    import torch
    from transformers import AutoTokenizer, AutoModelForSequenceClassification

    checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
    tokenizer = AutoTokenizer.from_pretrained(checkpoint)
    model = AutoModelForSequenceClassification.from_pretrained(checkpoint)

    sequence = "I've been waiting for a HuggingFace course my whole life."

    tokens = tokenizer.tokenize(sequence)
    ids = tokenizer.convert_tokens_to_ids(tokens)

    input_ids = torch.tensor([ids])
    print("Input IDs:", input_ids)

    output = model(input_ids)
    print("Logits:", output.logits)
```

```
⇒ Input IDs: tensor([[ 1045,  1005,  2310,  2042,  3403,  2005,  1037, 17662, 12172, 2607,
                    2026,  2878,  2166,  1012]])
    Logits: tensor([[ -2.7276,  2.8789]], grad_fn=<AddmmBackward0>)
```

Mengonversi urutan teks menjadi ID token dan kemudian mengemasnya ke dalam tensor PyTorch dengan menambahkan dimensi batch (menggunakan torch.tensor([ids])), yang diperlukan oleh model untuk memproses inputnya.

```
[ ] batched_ids = [ids, ids]

[ ] # Padding the inputs

    batched_ids = [
        [200, 200, 200],
        [200, 200]
    ]

[ ] padding_id = 100

    batched_ids = [
        [200, 200, 200],
        [200, 200, padding_id],
    ]

[ ] model = AutoModelForSequenceClassification.from_pretrained(checkpoint)

    sequence1_ids = [[200, 200, 200]]
    sequence2_ids = [[200, 200]]
    batched_ids = [
        [200, 200, 200],
        [200, 200, tokenizer.pad_token_id],
    ]

    print(model(torch.tensor(sequence1_ids)).logits)
    print(model(torch.tensor(sequence2_ids)).logits)
    print(model(torch.tensor(batched_ids)).logits)
```

⇒ We strongly recommend passing in an `attention_mask` since your input_ids may contain padding tokens.

```
tensor([[ 1.5694, -1.3895]], grad_fn=<AddmmBackward0>)
tensor([[ 0.5803, -0.4125]], grad_fn=<AddmmBackward0>)
tensor([[ 1.5694, -1.3895],
        [ 1.3374, -1.2163]], grad_fn=<AddmmBackward0>)
```

Memberikan beberapa jenis input ke model secara berurutan untuk melihat bagaimana model menangani tensor input yang berbeda

```
[ ] # Attention masks

batched_ids = [
    [200, 200, 200],
    [200, 200, tokenizer.pad_token_id],
]

[ ] attention_mask = [
    [1, 1, 1],
    [1, 1, 0],
]

outputs = model(torch.tensor(batched_ids), attention_mask=torch.tensor(attention_mask))
print(outputs.logits)
```

```
↳ tensor([[ 1.5694, -1.3895],
          [ 0.5803, -0.4125]], grad_fn=<AddmmBackward0>)
```

Pada kode ini, attention_mask digunakan untuk memberi tahu model token mana yang harus diperhatikan (1) dan token mana yang harus diabaikan (0), sehingga model hanya memproses token yang relevan dan mengabaikan padding saat menghitung logits.

```
[ ] #Longer sequences
sequence = sequence[:max_sequence_length]
```

```
↳ -----
NameError                                Traceback (most recent call last)
<ipython-input-18-e020d6cf7f6b> in <cell line: 2>()
      1 #Longer sequences
----> 2 sequence = sequence[:max_sequence_length]

NameError: name 'max_sequence_length' is not defined
```

```
[ ] #perbaikan :
# Tentukan panjang maksimal sequence
max_sequence_length = 128 # Misalnya, batas maksimal panjang sequence adalah 128 token

# Potong sequence sesuai panjang maksimal
sequence = sequence[:max_sequence_length]

print(sequence)
```

```
↳ [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
[ ] # Putting it all together
```

```
from transformers import AutoTokenizer
```

```
checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
```

```
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
```

```
sequence = "I've been waiting for a HuggingFace course my whole life."
```

```
model_inputs = tokenizer(sequence)
```

```
[ ] sequence = "I've been waiting for a HuggingFace course my whole life."
```

```
model_inputs = tokenizer(sequence)
```

```
[ ] sequences = ["I've been waiting for a HuggingFace course my whole life.", "So have I!"]
```

```
model_inputs = tokenizer(sequences)
```

```
[ ] # Will pad the sequences up to the maximum sequence length
```

```
model_inputs = tokenizer(sequences, padding="longest")
```

```
# Will pad the sequences up to the model max length
```

```
# (512 for BERT or DistilBERT)
```

```
model_inputs = tokenizer(sequences, padding="max_length")
```

```
# Will pad the sequences up to the specified max length
```

```
model_inputs = tokenizer(sequences, padding="max_length", max_length=8)
```

```
[ ] sequences = ["I've been waiting for a HuggingFace course my whole life.", "So have I!"]
```

```
# Will truncate the sequences that are longer than the model max length
```

```
# (512 for BERT or DistilBERT)
```

```
model_inputs = tokenizer(sequences, truncation=True)
```

```
# Will truncate the sequences that are longer than the specified max length
```

```
model_inputs = tokenizer(sequences, max_length=8, truncation=True)
```

```

sequences = ["I've been waiting for a HuggingFace course my whole life.", "So have I!"]

# Returns PyTorch tensors
model_inputs = tokenizer(sequences, padding=True, return_tensors="pt")

# Returns TensorFlow tensors
model_inputs = tokenizer(sequences, padding=True, return_tensors="tf")

# Returns NumPy arrays
model_inputs = tokenizer(sequences, padding=True, return_tensors="np")

[ ] # special tokens

sequence = "I've been waiting for a HuggingFace course my whole life."

model_inputs = tokenizer(sequence)
print(model_inputs["input_ids"])

tokens = tokenizer.tokenize(sequence)
ids = tokenizer.convert_tokens_to_ids(tokens)
print(ids)

⇒ [101, 1045, 1005, 2310, 2042, 3403, 2005, 1037, 17662, 12172, 2607, 2026, 2878, 2166, 1012, 102]
   [1045, 1005, 2310, 2042, 3403, 2005, 1037, 17662, 12172, 2607, 2026, 2878, 2166, 1012]

```

Pada kode ini, `tokenizer(sequence)` mengubah teks menjadi token dan input ID yang sesuai dengan model, termasuk token khusus seperti [CLS] dan [SEP] jika diperlukan oleh model. Sedangkan `tokenizer.tokenize(sequence)` memecah teks menjadi token-token individual dan `tokenizer.convert_tokens_to_ids(tokens)` mengonversi token-token tersebut menjadi ID yang bisa diproses oleh model.

```

[ ] print(tokenizer.decode(model_inputs["input_ids"]))
    print(tokenizer.decode(ids))

⇒ [CLS] i've been waiting for a huggingface course my whole life. [SEP]
   i've been waiting for a huggingface course my whole life.

```

Pada kode ini, `tokenizer.decode(model_inputs["input_ids"])` mengonversi ID input yang dihasilkan oleh tokenizer kembali ke bentuk teks asli, termasuk token khusus yang digunakan oleh model. Sedangkan `tokenizer.decode(ids)` mengonversi ID yang dihasilkan dari tokenisasi manual (tanpa mempertimbangkan token khusus) kembali menjadi teks, yang mungkin tidak mencakup beberapa token yang diperlukan untuk pemrosesan model.

```
[ ] # Wrapping up: From tokenizer to model

import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification

checkpoint = "distilbert-base-uncased-finetuned-sst-2-english"
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
model = AutoModelForSequenceClassification.from_pretrained(checkpoint)
sequences = ["I've been waiting for a HuggingFace course my whole life.", "So have I!"]

tokens = tokenizer(sequences, padding=True, truncation=True, return_tensors="pt")
output = model(**tokens)
```

```
[ ] # Processing the data
import torch
from transformers import AdamW, AutoTokenizer, AutoModelForSequenceClassification

# Same as before
checkpoint = "bert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
model = AutoModelForSequenceClassification.from_pretrained(checkpoint)
sequences = [
    "I've been waiting for a HuggingFace course my whole life.",
    "This course is amazing!",
]
batch = tokenizer(sequences, padding=True, truncation=True, return_tensors="pt")
```

```
tokenizer_config.json: 100% ██████████ 48.0/48.0 [00:00<00:00, 806B/s]
config.json: 100% ██████████ 570/570 [00:00<00:00, 11.5kB/s]
vocab.txt: 100% ██████████ 232k/232k [00:00<00:00, 2.03MB/s]
tokenizer.json: 100% ██████████ 466k/466k [00:00<00:00, 2.30MB/s]
model.safetensors: 100% ██████████ 440M/440M [00:04<00:00, 112MB/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
```

Kode ini mempersiapkan data untuk diproses menggunakan model BERT. Dengan menggunakan `AutoTokenizer.from_pretrained(checkpoint)`, tokenisasi dilakukan pada beberapa kalimat dalam `sequences`, menambahkan padding dan pemotongan agar ukuran input seragam. `return_tensors="pt"` menghasilkan output dalam format tensor PyTorch, yang siap untuk diberikan ke model BERT (`AutoModelForSequenceClassification.from_pretrained(checkpoint)`) untuk klasifikasi urutan.

```
[ ] # This is new
batch["labels"] = torch.tensor([1, 1])

optimizer = AdamW(model.parameters())
loss = model(**batch).loss
loss.backward()
optimizer.step()
```

```
/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:591: FutureWarning: This implementation of AdamW is deprecated. Please use the implementation in torch.nn.optim.AdamW instead.
warnings.warn(
```

Kode ini menunjukkan bagaimana melakukan pelatihan dengan model BERT. Setelah menambahkan label ke dalam batch (dalam hal ini, dua label yang bernilai 1), optimizer `AdamW` digunakan untuk memperbarui parameter model. Langkah-langkahnya adalah sebagai berikut:

1. Menghitung loss dengan memberikan batch ke model.
2. Melakukan backward pass untuk menghitung gradien (`loss.backward()`).
3. Menggunakan optimizer untuk memperbarui parameter model berdasarkan gradien yang dihitung (`optimizer.step()`).

Proses ini merupakan bagian dari training loop untuk mengoptimalkan model.

Masalah:

```
[ ] #Loading a dataset from the Hub

from datasets import load_dataset

raw_datasets = load_dataset("glue", "mrpc")
raw_datasets
DatasetDict({
  train: Dataset({
    features: ['sentence1', 'sentence2', 'label', 'idx'],
    num_rows: 3668
  })
  validation: Dataset({
    features: ['sentence1', 'sentence2', 'label', 'idx'],
    num_rows: 408
  })
  test: Dataset({
    features: ['sentence1', 'sentence2', 'label', 'idx'],
    num_rows: 1725
  })
})
```

```
File "<ipython-input-33-905b91105edf>", line 8
  train: Dataset({
    ^
SyntaxError: invalid syntax. Perhaps you forgot a comma?
```

Solusi:



✓ Connected to Python 3 Google Compute Engine backend


```

[11] raw_train_dataset = raw_datasets["train"]
    raw_train_dataset[0]

⇒ {'sentence1': 'Amrozi accused his brother , whom he called " the witness " , of deliberately distorting his
    evidence .',
    'sentence2': 'Referring to him as only " the witness " , Amrozi accused his brother of deliberately
    distorting his evidence .',
    'label': 1,
    'idx': 0}

```

Kode tersebut mengakses dataset pelatihan (train dataset) dari sebuah objek raw_datasets dan menampilkan data pertama pada dataset tersebut. raw_datasets["train"] adalah bagian dataset yang berisi data pelatihan, dan raw_train_dataset[0] mengakses item pertama dari dataset pelatihan.

Jika raw_datasets adalah dictionary yang berisi dataset terstruktur, maka raw_train_dataset[0] kemungkinan akan mengembalikan data yang berisi teks dan label atau fitur lainnya yang ada dalam dataset pelatihan tersebut.

```

[15] inputs = tokenizer("This is the first sentence.", "This is the second one.")
    inputs

⇒ {'input_ids': [101, 2023, 2003, 1996, 2034, 6251, 1012, 102, 2023, 2003, 1996, 2117, 2028, 1012, 102],
    'token_type_ids': [0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1], 'attention_mask': [1, 1, 1, 1, 1, 1, 1, 1,
    1, 1, 1, 1, 1, 1, 1]}

```

Kode tersebut menggunakan tokenizer untuk mengonversi dua kalimat ("This is the first sentence." dan "This is the second one.") menjadi format yang dapat diproses oleh model.

```

[16] tokenizer.convert_ids_to_tokens(inputs["input_ids"])

⇒ ['[CLS]',
    'this',
    'is',
    'the',
    'first',
    'sentence',
    '.',
    '[SEP]',
    'this',
    'is',
    'the',
    'second',
    'one',
    '.',
    '[SEP]']

```

Fungsi tokenizer.convert_ids_to_tokens(inputs["input_ids"]) digunakan untuk mengonversi ID token kembali ke kata-kata atau token yang dapat dibaca manusia. inputs["input_ids"] berisi

ID token yang telah dihasilkan sebelumnya, dan fungsi ini akan mengembalikan daftar kata atau token yang sesuai dengan ID tersebut.

```
[17] tokenized_dataset = tokenizer(  
    raw_datasets["train"]["sentence1"],  
    raw_datasets["train"]["sentence2"],  
    padding=True,  
    truncation=True,  
)
```

```
def tokenize_function(example):  
    return tokenizer(example["sentence1"], example["sentence2"], truncation=True)
```

```
[19] tokenized_datasets = raw_datasets.map(tokenize_function, batched=True)  
tokenized_datasets
```

Map: 100% 3668/3668 [00:01<00:00, 2289.11 examples/s]
Map: 100% 408/408 [00:00<00:00, 1986.80 examples/s]
Map: 100% 1725/1725 [00:00<00:00, 3104.93 examples/s]

```
DatasetDict({  
  train: Dataset({  
    features: ['sentence1', 'sentence2', 'label', 'idx', 'input_ids', 'token_type_ids',  
    'attention_mask'],  
    num_rows: 3668  
  })  
  validation: Dataset({  
    features: ['sentence1', 'sentence2', 'label', 'idx', 'input_ids', 'token_type_ids',  
    'attention_mask'],  
    num_rows: 408  
  })  
  test: Dataset({  
    features: ['sentence1', 'sentence2', 'label', 'idx', 'input_ids', 'token_type_ids',  
    'attention_mask'],  
    num_rows: 1725  
  })  
})
```

tokenized_datasets adalah hasil dari pemetaan fungsi tokenisasi pada dataset mentah (raw_datasets) untuk mengubah data teks menjadi format yang dapat diproses oleh model pembelajaran mesin.

Kendala:

```
[20] DatasetDict({
  train: Dataset({
    features: ['attention_mask', 'idx', 'input_ids', 'label', 'sent
    num_rows: 3668
  })
  validation: Dataset({
    features: ['attention_mask', 'idx', 'input_ids', 'label', 'sent
    num_rows: 408
  })
  test: Dataset({
    features: ['attention_mask', 'idx', 'input_ids', 'label', 'sent
    num_rows: 1725
  })
})

File "<ipython-input-20-e65972222e0a>", line 2
  train: Dataset({
    ^
SyntaxError: invalid syntax. Perhaps you forgot a comma?
```

Error tersebut terjadi karena format DatasetDict yang Anda masukkan ke interpreter Python tidak valid secara sintaksis. Dalam Python, dictionary didefinisikan menggunakan pasangan kunci dan nilai yang dipisahkan oleh tanda titik dua (:), dan setiap pasangan dipisahkan oleh tanda koma (,). Anda juga tidak perlu menuliskan data dalam format literal Python seperti output log.

Solusi:

```
✓ [22] from datasets import load_dataset
2s

# Memuat dataset GLUE (MRPC)
raw_datasets = load_dataset("glue", "mrpc")

# Menampilkan struktur dataset
print("Dataset berhasil dimuat:")
print(raw_datasets)

Dataset berhasil dimuat:
DatasetDict({
  train: Dataset({
    features: ['sentence1', 'sentence2', 'label', 'idx'],
    num_rows: 3668
  })
  validation: Dataset({
    features: ['sentence1', 'sentence2', 'label', 'idx'],
    num_rows: 408
  })
  test: Dataset({
    features: ['sentence1', 'sentence2', 'label', 'idx'],
    num_rows: 1725
  })
})
```

```
✓ [23] # Dynamic padding
0s
from transformers import DataCollatorWithPadding

data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

samples = tokenized_datasets["train"][:8]
samples = {k: v for k, v in samples.items() if k not in ["idx", "sentence1", "sentence2"]}
[ len(x) for x in samples["input_ids"] ]

[50, 59, 47, 67, 59, 50, 62, 32]
```

Kode tersebut menggunakan `DataCollatorWithPadding` untuk menambahkan padding secara dinamis pada tokenisasi data, memastikan bahwa setiap sampel dalam batch memiliki panjang yang konsisten

```
[24] batch = data_collator(samples)
      {k: v.shape for k, v in batch.items()}
```

```
⇒ {'input_ids': torch.Size([8, 67]),
   'token_type_ids': torch.Size([8, 67]),
   'attention_mask': torch.Size([8, 67]),
   'labels': torch.Size([8])}
```

Kode tersebut menggunakan data_collator untuk menggabungkan sampel menjadi satu batch dan kemudian memeriksa dimensi dari setiap item dalam batch tersebut.

```
✓ [25] # Fine-tuning a model with the Trainer API
3s
      from datasets import load_dataset
      from transformers import AutoTokenizer, DataCollatorWithPadding

      raw_datasets = load_dataset("glue", "mrpc")
      checkpoint = "bert-base-uncased"
      tokenizer = AutoTokenizer.from_pretrained(checkpoint)
```

```
✓ [26] def tokenize_function(example):
1s
        return tokenizer(example["sentence1"], example["sentence2"], truncation=True)

      tokenized_datasets = raw_datasets.map(tokenize_function, batched=True)
      data_collator = DataCollatorWithPadding(tokenizer=tokenizer)
```

```
⇒ Map: 100% ██████████ 3668/3668 [00:00<00:00, 4113.97 examples/s]
```

memuat dan menyiapkan dataset GLUE MRPC untuk fine-tuning model menggunakan API Trainer dari Hugging Face.

```
✓ [27] # Training
0s
      from transformers import TrainingArguments

      training_args = TrainingArguments("test-trainer")
```

```
✓ [28] from transformers import AutoModelForSequenceClassification
9s
      model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)
```

```
⇒ model.safetensors: 100% ██████████ 440M/440M [00:07<00:00, 67.7MB/s]
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and

berfokus pada inisialisasi argumen pelatihan dan model untuk fine-tuning menggunakan API Trainer dari Hugging Face.

```
[29] from transformers import Trainer

trainer = Trainer(
    model,
    training_args,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["validation"],
    data_collator=data_collator,
    tokenizer=tokenizer,
)

<ipython-input-29-36fd205b194c>:3: FutureWarning: `tokenizer` is deprecated and will be removed in version 5
trainer = Trainer(
```

Kode ini menyusun objek Trainer yang digunakan untuk melatih model dengan dataset yang telah diproses.

```
[30] trainer.train()

wandb: WARNING The `run_name` is currently set to the same value as `TrainingArguments.output_dir`. If this
wandb: Using wandb-core as the SDK backend. Please refer to https://wandb.me/wandb-core for more informati
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wandb.me/wandb-server)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit: .....
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
Tracking run with wandb version 0.19.1
Run data is saved locally in /content/wandb/run-20250104_090312-iwjav9hx
Syncing run test-trainer to Weights & Biases (docs)
View project at https://wandb.ai/rizkiapriliarahman/huggingface
View run at https://wandb.ai/rizkiapriliarahman/huggingface/runs/iwjav9hx
[ 323/1377 37:42 < 2:03:48, 0.14 it/s, Epoch 0.70/3]

Step Training Loss
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-30-3435b262f1ae> in <cell line: 1>()
----> 1 trainer.train()

6 frames
/usr/local/lib/python3.10/dist-packages/torch/autograd/graph.py in _engine_run_backward(t_outputs, *args,
```

Kendala: Running sangat lama dan selalu error karena data sangat banyak.

```
[31] # Evaluation

predictions = trainer.predict(tokenized_datasets["validation"])
print(predictions.predictions.shape, predictions.label_ids.shape)

[ 323/1377 37:42 < 2:03:48, 0.14 it/s, Epoch 0.70/3]

Step Training Loss
-----
[51/51 02:09]
(408, 2) (408,)
```

Kode ini digunakan untuk mengevaluasi model yang telah dilatih dengan menggunakan dataset validasi dan melihat hasil prediksi serta label yang terkait.

```
✓ [32] import numpy as np

      preds = np.argmax(predictions.predictions, axis=-1)
```

Kode ini digunakan untuk mengonversi hasil prediksi model yang berupa probabilitas atau skor ke dalam bentuk label kelas yang dipilih (klasifikasi argmax).

Kendala:

```
⌚ 0s ▶ import evaluate

metric = evaluate.load("glue", "mrpc")
metric.compute(predictions=preds, references=predictions.label_ids)
```

⇌

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
<ipython-input-33-079868121e9c> in <cell line: 1>()
----> 1 import evaluate
      2
      3 metric = evaluate.load("glue", "mrpc")
      4 metric.compute(predictions=preds, references=predictions.label_ids)

ModuleNotFoundError: No module named 'evaluate'

-----

NOTE: If your import is failing due to a missing package, you can
manually install dependencies using either !pip or !apt.

To view examples of installing some common dependencies, click the
"Open Examples" button below.

-----

OPEN EXAMPLES
```

Next steps: [Explain error](#)

Solusi:

```
✓ 11s pip install evaluate

Collecting evaluate
  Downloading evaluate-0.4.3-py3-none-any.whl.metadata (9.2 kB)
Requirement already satisfied: datasets>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.20.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.24.3)
Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.3.8)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.2.2)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.31.0)
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from evaluate) (4.66.1)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from evaluate) (3.5.0)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from evaluate) (3.9.1)
Requirement already satisfied: fsspec>=2021.05.0 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]>=2021.05.0->evaluate) (2024.9.0)
Requirement already satisfied: huggingface-hub>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.24.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from evaluate) (24.2)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (3.16.1)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (15.0.2)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (3.10.10)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (6.0.2)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (4.12.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate) (2024.12.14)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->evaluate) (2.9.0)

✓ 2s [35] import evaluate

metric = evaluate.load("glue", "mrpc")
metric.compute(predictions=preds, references=predictions.label_ids)

Downloading builder script: 100% 5.75k/5.75k [00:00<00:00, 9.46kB/s]
{'accuracy': 0.8112745098039216, 'f1': 0.851063829787234}
```

Kode ini digunakan untuk mengevaluasi kinerja model pada tugas MRPC menggunakan metrik dari GLUE (General Language Understanding Evaluation).

```

✓ [36] def compute_metrics(eval_preds):
0s      metric = evaluate.load("glue", "mrpc")
      logits, labels = eval_preds
      predictions = np.argmax(logits, axis=-1)
      return metric.compute(predictions=predictions, references=labels)

✓ [37] training_args = TrainingArguments("test-trainer", evaluation_strategy="epoch")
1s      model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)

      trainer = Trainer(
          model,
          training_args,
          train_dataset=tokenized_datasets["train"],
          eval_dataset=tokenized_datasets["validation"],
          data_collator=data_collator,
          tokenizer=tokenizer,
          compute_metrics=compute_metrics,
      )

➡ /usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1575: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.0.0. Please use `eval_strategy` instead.
warnings.warn(
    Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
)
<ipython-input-37-415be425bb02>:4: FutureWarning: `tokenizer` is deprecated and will be removed in version 4.0.0. Please use `tokenizer_name` instead.
trainer = Trainer(

```

Kode ini menyusun kembali objek Trainer dengan menambahkan beberapa parameter untuk kontrol lebih lanjut atas proses pelatihan dan evaluasi model.

```

⏱ 12m [38] trainer.train()

➡ [ 101/1377 12:30 < 2:41:16, 0.13 it/s, Epoch 0.22/3]

Epoch Training Loss Validation Loss
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-38-3435b262f1ae> in <cell line: 1>()
----> 1 trainer.train()

^ 5 frames

```

Kendala: Running sangat lama karena membuat data yang banyak.

```

✓ 3s [39] from datasets import load_dataset
      from transformers import AutoTokenizer, DataCollatorWithPadding

      raw_datasets = load_dataset("glue", "mrpc")
      checkpoint = "bert-base-uncased"
      tokenizer = AutoTokenizer.from_pretrained(checkpoint)

✓ 0s [40] def tokenize_function(example):
      return tokenizer(example["sentence1"], example["sentence2"], truncation=True)

      tokenized_datasets = raw_datasets.map(tokenize_function, batched=True)
      data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

➡ Map: 100% 408/408 [00:00<00:00, 2418.17 examples/s]

```

Kode ini digunakan untuk menyiapkan dataset dan tokenizer sebelum melatih model.

```

✓ 0s [41] # Prepare for training

      tokenized_datasets = tokenized_datasets.remove_columns(["sentence1", "sentence2", "idx"])
      tokenized_datasets = tokenized_datasets.rename_column("label", "labels")
      tokenized_datasets.set_format("torch")
      tokenized_datasets["train"].column_names

➡ ['labels', 'input_ids', 'token_type_ids', 'attention_mask']

```

Kode ini mempersiapkan dataset yang telah ditokenisasi agar dapat digunakan dalam pelatihan model.

```

✓ 0s [42] from torch.utils.data import DataLoader

      train_dataloader = DataLoader(
          tokenized_datasets["train"], shuffle=True, batch_size=8, collate_fn=data_collator
      )
      eval_dataloader = DataLoader(
          tokenized_datasets["validation"], batch_size=8, collate_fn=data_collator
      )

✓ 1s [43] for batch in train_dataloader:
      break
      {k: v.shape for k, v in batch.items()}

➡ {'labels': torch.Size([8]),
   'input_ids': torch.Size([8, 69]),
   'token_type_ids': torch.Size([8, 69]),
   'attention_mask': torch.Size([8, 69])}

```

Kode ini digunakan untuk melihat bentuk (shape) dari data dalam batch pertama yang diambil dari train_dataloader.

```
[44] from transformers import AutoModelForSequenceClassification

model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)

outputs = model(**batch)
print(outputs.loss, outputs.logits.shape)
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint. You should probably TRAIN this model on a down-stream task to be able to use it for predictions.

Kode ini digunakan untuk melakukan inferensi atau pelatihan pada model dengan batch data yang telah dipersiapkan.

```
[45] from transformers import AdamW

optimizer = AdamW(model.parameters(), lr=5e-5)
```

/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:591: FutureWarning: The `warn` method is deprecated. Please use `warnings.warn` instead.

Kode ini digunakan untuk mengonfigurasi optimizer yang akan digunakan selama pelatihan model.

```
[46] from transformers import get_scheduler

num_epochs = 3
num_training_steps = num_epochs * len(train_dataloader)
lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps,
)
print(num_training_steps)
```

1377

Kode ini digunakan untuk menyiapkan scheduler yang akan mengatur laju pembelajaran (learning rate) selama pelatihan.

```
[48] # The training loop

import torch

device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
model.to(device)
device

⇒ device(type='cpu')
```

Kode ini digunakan untuk mempersiapkan perangkat (device) tempat model akan dilatih, apakah itu GPU (CUDA) atau CPU.

```
from tqdm.auto import tqdm

progress_bar = tqdm(range(num_training_steps))

model.train()
for epoch in range(num_epochs):
    for batch in train_dataloader:
        batch = {k: v.to(device) for k, v in batch.items()}
        outputs = model(**batch)
        loss = outputs.loss
        loss.backward()

        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()
        progress_bar.update(1)
```

⇒ 26% 357/1377 [2:08:16<2:02:25, 7.20s/it]

```
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-49-b86b10a2f8e7> in <cell line: 6>()
      9         outputs = model(**batch)
     10         loss = outputs.loss
--> 11         loss.backward()
     12
```

Kendala: Running sangat lama dan error dikarenakan data yg sangat banyak.

✓
1m

[50] # The evaluation loop

```
import evaluate

metric = evaluate.load("glue", "mrpc")
model.eval()
for batch in eval_dataloader:
    batch = {k: v.to(device) for k, v in batch.items()}
    with torch.no_grad():
        outputs = model(**batch)

    logits = outputs.logits
    predictions = torch.argmax(logits, dim=-1)
    metric.add_batch(predictions=predictions, references=batch["labels"])

metric.compute()
```

⇒ {'accuracy': 0.8406862745098039, 'f1': 0.8881239242685026}

Kode ini digunakan untuk melakukan evaluasi model setelah pelatihan.

```
#Supercharge your training loop with 🚀 Accelerate

from transformers import AdamW, AutoModelForSequenceClassification, get_scheduler

model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)
optimizer = AdamW(model.parameters(), lr=3e-5)

device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
model.to(device)

num_epochs = 3
num_training_steps = num_epochs * len(train_dataloader)
lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps,
)

progress_bar = tqdm(range(num_training_steps))


model.train()
for epoch in range(num_epochs):
    for batch in train_dataloader:
        batch = {k: v.to(device) for k, v in batch.items()}
        outputs = model(**batch)
        loss = outputs.loss
        loss.backward()

        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()
        progress_bar.update(1)
```

⚠️ The weights of BertForSequenceClassification were not initialized from the model checkpoint. You should probably TRAIN this model on a down-stream task to be able to use it for predictions.

sr/local/lib/python3.10/dist-packages/transformers/optimization.py:591: FutureWarning: T

warnings.warn(

5%  201/1377 [1:23:15<2:10:37, 6.66s/it]

/boardInterrupt Traceback (most recent call last)

python-innnt-51-1069be20df2e> in <cell line: 23>()

File ~/... completed at 0:16 PM

Kendala: Running sangat lama karena data yang sangat banyak.

Kendala:

```
[60] # And here are the changes:

from transformers import AdamW, AutoModelForSequenceClassification, get_scheduler

model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)
optimizer = AdamW(model.parameters(), lr=3e-5)

device = torch.device("cpu") if torch.cuda.is_available() else torch.device("cpu")
model.to(device)

num_epochs = 3
num_training_steps = num_epochs * len(train_dataloader)
lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps,
)

progress_bar = tqdm(range(num_training_steps))

model.train()
for epoch in range(num_epochs):
    for batch in train_dataloader:
        batch = {k: v.to(device) for k, v in batch.items()}
        outputs = model(**batch)
        loss = outputs.loss
        loss.backward()

        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()
        progress_bar.update(1)

-----
HTTPError                                Traceback (most recent call last)
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_http.py in hf_raise_for_status(response,
endpoint_name)
```

Solusi:

```
[61] pip install --upgrade transformers

Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.34.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (3.12.2)
Requirement already satisfied: huggingface-hub<0.17.0,>=0.16.0 in /usr/local/lib/python3.10/dist-packages (0.16.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (1.24.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (23.1)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (6.0.1)
Requirement already satisfied: regex in /usr/local/lib/python3.10/dist-packages (2023.10.3)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (2.31.0)
Requirement already satisfied: tokenizers in /usr/local/lib/python3.10/dist-packages (0.15.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.65.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (4.5.0)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (2.0.3)
```




And here are the changes:

```
from transformers import AdamW, AutoModelForSequenceClassification, get_s
checkpoint = "bert-base-uncased" # Nama model dari Hugging Face
# atau jika menggunakan model lokal
checkpoint = "C:/Users/Asus/Downloads/Semester 7/ML"

device = torch.device("cpu") if torch.cuda.is_available() else torch.device("cpu")
model.to(device)

num_epochs = 3
num_training_steps = num_epochs * len(train_dataloader)
lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps,
)

progress_bar = tqdm(range(num_training_steps))

model.train()
for epoch in range(num_epochs):
    for batch in train_dataloader:
        batch = {k: v.to(device) for k, v in batch.items()}
        outputs = model(**batch)
        loss = outputs.loss
        loss.backward()

        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()
        progress_bar.update(1)
```



16%



222/1377 [45:19<3:54:45, 12.20s/it]

KeyboardInterrupt

Traceback (most recent call last):

<ipython-input-64-66ec70968799> in <cell line: 23>()

26 outputs = model(**batch)

27 loss = outputs.loss

! 0s completed at 9:16 PM

Kendala:

```
from accelerate import Accelerator
from transformers import AdamW, AutoModelForSequenceClassification, get_scheduler

accelerator = Accelerator()

model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)
optimizer = AdamW(model.parameters(), lr=3e-5)

train_dl, eval_dl, model, optimizer = accelerator.prepare(
    train_dataloader, eval_dataloader, model, optimizer
)

num_epochs = 3
num_training_steps = num_epochs * len(train_dl)
lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps,
)

progress_bar = tqdm(range(num_training_steps))

model.train()
for epoch in range(num_epochs):
    for batch in train_dl:
        outputs = model(**batch)
        loss = outputs.loss
        accelerator.backward(loss)

        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()
        progress_bar.update(1)

-----
HTTPError                                Traceback (most recent call last)
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_http.py in hf_raise_for_status(response)
    405     try:
--> 406         response.raise_for_status()
```

Variabel checkpoint Tidak Didefinisikan Tambahkan definisi untuk variabel checkpoint dengan model yang valid. Data Tidak Didefinisikan train_dataloader dan eval_dataloader perlu didefinisikan sebelumnya. batch Tidak Dipindahkan ke Perangkat Perlu memindahkan batch ke perangkat (device) jika tidak menggunakan Accelerator untuk itu. tqdm Tidak Diimpor Tambahkan impor tqdm jika belum dilakukan.

Solusi:

```
▶ from accelerate import Accelerator
from transformers import AdamW, AutoModelForSequenceClassification, get_scheduler
from tqdm import tqdm
import torch

# Definisi checkpoint model (sesuaikan dengan model yang ingin digunakan)
checkpoint = "bert-base-uncased"

# Inisialisasi Accelerator
accelerator = Accelerator()

# Memuat model
model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)

# Optimizer
optimizer = AdamW(model.parameters(), lr=3e-5)

# Asumsi train_dataloader dan eval_dataloader sudah dibuat sebelumnya
# Contoh sederhana untuk memeriksa data
# from transformers import AutoTokenizer
# tokenizer = AutoTokenizer.from_pretrained(checkpoint)
# train_dataset = ...
# train_dataloader = DataLoader(train_dataset, batch_size=8)

# Mempersiapkan komponen dengan Accelerator
train_dl, eval_dl, model, optimizer = accelerator.prepare(
    train_dataloader, eval_dataloader, model, optimizer
)

# Scheduler untuk pembelajaran
num_epochs = 3
num_training_steps = num_epochs * len(train_dl)
lr_scheduler = get_scheduler(
    "linear",
    optimizer=optimizer,
    num_warmup_steps=0,
    num_training_steps=num_training_steps,
)

# Progress bar
progress_bar = tqdm(range(num_training_steps))
```

```

# Melatih model
model.train()
for epoch in range(num_epochs):
    for batch in train_dl:
        # Memindahkan batch ke perangkat (dengan Accelerator ini otomatis)
        outputs = model(**batch)
        loss = outputs.loss

        # Backpropagation dengan Accelerator
        accelerator.backward(loss)

        # Optimizer step
        optimizer.step()
        lr_scheduler.step()
        optimizer.zero_grad()

        # Perbarui progress bar
        progress_bar.update(1)

print("Training selesai!")

```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint. You should probably TRAIN this model on a down-stream task to be able to use it.

0%| | 0/1377 [00:00<?, ?it/s]

0%| | 1/1377 [00:12<4:41:39, 12.28s/it]

0%| | 2/1377 [00:20<3:52:16, 10.14s/it]

0%| | 3/1377 [00:29<3:37:25, 9.49s/it]

0%| | 4/1377 [00:40<3:50:52, 10.09s/it]

0%| | 5/1377 [00:46<3:12:12, 8.41s/it]

checkpoint: Pastikan model yang digunakan valid, seperti bert-base-uncased.

DataLoader: perlu memastikan bahwa train_dataloader dan eval_dataloader sudah didefinisikan dengan dataset yang benar.

`accelerator.backward()`: Metode ini menggantikan `loss.backward()` untuk mendukung multi-GPU atau mixed precision.

Progress Bar: Progress bar diatur menggunakan `tqdm` untuk memantau langkah selama pelatihan.

menambahkan ! agar bisa di running bukan di terminal

```
✓ [69] !accelerate launch train.py
24s
➔ The following values were not passed to `accelerate launch` and had defaults used in
  `--num_processes` was set to a value of `0`
  `--num_machines` was set to a value of `1`
  `--mixed_precision` was set to a value of `no`
  `--dynamo_backend` was set to a value of `no`
To avoid this warning pass in values for each of the problematic parameters or run `
/usr/bin/python3: can't open file '/content/train.py': [Errno 2] No such file or dir
Traceback (most recent call last):
  File "/usr/local/bin/accelerate", line 8, in <module>
```

digunakan untuk meluncurkan pelatihan model menggunakan *Accelerate*, pustaka dari Hugging Face yang dirancang untuk menyederhanakan proses pelatihan model di berbagai perangkat (misalnya, satu atau lebih GPU, atau CPU) dengan konfigurasi yang lebih mudah dan lebih efisien.

```
✓ 11s # The Hugging Face Hub

#Using pretrained models
from transformers import pipeline

camembert_fill_mask = pipeline("fill-mask", model="camembert-base")
results = camembert_fill_mask("Le camembert est <mask> :)")

➔ config.json: 100% ██████████ 508/508 [00:00<00:00, 7.37kB/s]
model.safetensors: 100% ██████████ 445M/445M [00:06<00:00, 91.5MB/s]
Some weights of the model checkpoint at camembert-base were not used when initializing CamembertForMaskedLM
- This IS expected if you are initializing CamembertForMaskedLM from the checkpoint of a model trained on a
- This IS NOT expected if you are initializing CamembertForMaskedLM from the checkpoint of a model that you
tokenizer_config.json: 100% ██████████ 25.0/25.0 [00:00<00:00, 654B/s]
sentencepiece.bpe.model: 100% ██████████ 811k/811k [00:00<00:00, 3.25MB/s]
tokenizer.json: 100% ██████████ 1.40M/1.40M [00:00<00:00, 4.29MB/s]
Device set to use cpu
```

Kode ini menggunakan pipeline dari Hugging Face untuk melakukan tugas pengisian masker pada model bahasa *Camembert*.

```
[79] [
  {'sequence': 'Le camembert est délicieux :)', 'score': 0.49091005325317383, 'token': 7200, 'token_str': 'é'},
  {'sequence': 'Le camembert est excellent :)', 'score': 0.1055697426199913, 'token': 2183, 'token_str': 'e'},
  {'sequence': 'Le camembert est succulent :)', 'score': 0.03453313186764717, 'token': 26202, 'token_str': 's'},
  {'sequence': 'Le camembert est meilleur :)', 'score': 0.0330314114689827, 'token': 528, 'token_str': 'mei'},
  {'sequence': 'Le camembert est parfait :)', 'score': 0.03007650189101696, 'token': 1654, 'token_str': 'pa'}
]

[{'sequence': 'Le camembert est délicieux :)',
 'score': 0.49091005325317383,
 'token': 7200,
 'token_str': 'délicieux'},
 {'sequence': 'Le camembert est excellent :)',
 'score': 0.1055697426199913,
 'token': 2183,
 'token_str': 'excellent'},
 {'sequence': 'Le camembert est succulent :)',
 'score': 0.03453313186764717,
 'token': 26202,
 'token_str': 'succulent'},
 {'sequence': 'Le camembert est meilleur :)',
 'score': 0.0330314114689827,
 'token': 528,
 'token_str': 'meilleur'},
 {'sequence': 'Le camembert est parfait :)',
 'score': 0.03007650189101696,
 'token': 1654,
 'token_str': 'parfait'}]]
```

berisi beberapa prediksi dari model *Camembert* untuk mengisi masker dalam kalimat "Le camembert est <mask> :)"

```
[80] from transformers import CamembertTokenizer, CamembertForMaskedLM

tokenizer = CamembertTokenizer.from_pretrained("camembert-base")
model = CamembertForMaskedLM.from_pretrained("camembert-base")

Some weights of the model checkpoint at camembert-base were not used when initializing CamembertForMaskedLM:
- This IS expected if you are initializing CamembertForMaskedLM from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPretraining checkpoint).
- This IS NOT expected if you are initializing CamembertForMaskedLM from the checkpoint of a model that you expect to be initialized from another checkpoint (e.g. initializing a BertForPretraining model from a BertForPretraining checkpoint).
```

Kode ini memuat model dan tokenizer *Camembert* untuk tugas pengisian masker.

```
from transformers import AutoTokenizer, AutoModelForMaskedLM

tokenizer = AutoTokenizer.from_pretrained("camembert-base")
model = AutoModelForMaskedLM.from_pretrained("camembert-base")

Some weights of the model checkpoint at camembert-base were not used when initializing CamembertForMaskedLM:
- This IS expected if you are initializing CamembertForMaskedLM from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPretraining checkpoint).
- This IS NOT expected if you are initializing CamembertForMaskedLM from the checkpoint of a model that you expect to be initialized from another checkpoint (e.g. initializing a BertForPretraining model from a BertForPretraining checkpoint).
```

Kode ini menggunakan *AutoTokenizer* dan *AutoModelForMaskedLM* dari pustaka *transformers* untuk memuat model dan tokenizer *Camembert* dengan cara yang lebih generik, yang memungkinkan Anda menggunakan model dan tokenizer yang dapat dipertukarkan antara berbagai jenis model tanpa menyebutkan model secara spesifik.

```
[90] # Using the push_to_hub API

from huggingface_hub import notebook_login

notebook_login()

e properly logged in by executing 'huggingface-cli login', and if you did pa

[91] from transformers import TrainingArguments

training_args = TrainingArguments(
    "bert-finetuned-mrpc", save_strategy="epoch", push_to_hub=True
)

[92] from transformers import AutoModelForMaskedLM, AutoTokenizer

checkpoint = "camembert-base"

model = AutoModelForMaskedLM.from_pretrained(checkpoint)
tokenizer = AutoTokenizer.from_pretrained(checkpoint)

Some weights of the model checkpoint at camembert-base were not used when initializing CamembertForMaskedLM:
- This IS expected if you are initializing CamembertForMaskedLM from the checkpoint of a model trained on an
- This IS NOT expected if you are initializing CamembertForMaskedLM from the checkpoint of a model that you
```

```
[ ] !pip install datasets evaluate transformers[sentencepiece]
!apt install git-lfs

Collecting datasets
  Downloading datasets-3.2.0-py3-none-any.whl.metadata (20 kB)
Collecting evaluate
  Downloading evaluate-0.4.3-py3-none-any.whl.metadata (9.2 kB)
Requirement already satisfied: transformers[sentencepiece] in /usr/local/lib/python3.10.12/dist-packages (from datasets)
Requirement already satisfied: filelock in /usr/local/lib/python3.10.12/dist-packages (from datasets)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10.12/dist-packages (from datasets)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10.12/dist-packages (from datasets)
Collecting dill<0.3.9,>=0.3.0 (from datasets)
  Downloading dill-0.3.8-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: pandas in /usr/local/lib/python3.10.12/dist-packages (from datasets)
```

Dengan menjalankan kedua perintah tersebut, akan mendapatkan pustaka yang diperlukan untuk bekerja dengan dataset dan model dari Hugging Face, serta kemampuan untuk mengelola file besar melalui Git LFS.

```
[ ] !git config --global user.email "rizkiapriarahman1603@gmail.com"
    !git config --global user.name "rizkiapriarahman"
```

You will also need to be logged in to the Hugging Face Hub. Execute the following with your credentials.

need to setup git. So, I adapt my email and name in the following cell.

```
[ ] from huggingface_hub import notebook_login

    notebook_login()
```

Perintah `notebook_login()` dari pustaka `huggingface_hub` digunakan untuk masuk ke akun Hugging Face Hub dalam lingkungan notebook.

```
[ ] from transformers import TrainingArguments

    training_args = TrainingArguments(
        "bert-finetuned-mrpc", save_strategy="epoch", push_to_hub=True
    )
```

Perintah ini mengonfigurasi `TrainingArguments` untuk melatih model menggunakan *Transformers* library

```
▶ from transformers import AutoModelForMaskedLM, AutoTokenizer

    checkpoint = "camembert-base"

    model = AutoModelForMaskedLM.from_pretrained(checkpoint)
    tokenizer = AutoTokenizer.from_pretrained(checkpoint)
```

```
⇒ Some weights of the model checkpoint at camembert-base were not used
- This IS expected if you are initializing CamembertForMaskedLM from
- This IS NOT expected if you are initializing CamembertForMaskedLM
```

Kode ini digunakan untuk memuat model *Camembert* dan tokenizer-nya, yang dapat digunakan untuk tugas *Masked Language Modeling* (MLM),


```
[ ] tokenizer.push_to_hub("rizkiapriliarahman")
```

```
⇒ sentencepiece.bpe.model: 100% ██████████ 811k/811k [00:00<
CommitInfo(commit_url='https://huggingface.co/rizkiapriliarahman/riz
commit_message='Upload tokenizer', commit_description='',
oid='02930cae050c0d4add6b296c77d9e0d89f3b7084', pr_url=None,
repo_url=RepoUrl('https://huggingface.co/rizkiapriliarahman/rizkiapr
endpoint='https://huggingface.co', repo_type='model',
repo_id='rizkiapriliarahman/rizkiapriliarahman'), pr_revision=None,
```

Perintah `tokenizer.push_to_hub("rizkiapriliarahman")` digunakan untuk mengunggah tokenizer ke Hugging Face Hub dengan nama repositori **rizkiapriliarahman**. Setelah berhasil diunggah, tokenizer akan tersedia di Hugging Face Hub dan dapat diakses atau digunakan oleh siapa saja dengan repositori tersebut.

```
[ ] from huggingface_hub import create_repo
```

```
create_repo("tubes-ml-apil")
```

```
⇒ RepoUrl('https://huggingface.co/rizkiapriliarahman/tubes-ml-apil',
endpoint='https://huggingface.co', repo_type='model',
repo_id='rizkiapriliarahman/tubes-ml-apil')
```

Perintah `create_repo("tubes-ml-apil")` digunakan untuk membuat repositori baru di Hugging Face Hub dengan nama **tubes-ml-apil**.

```
[ ] from huggingface_hub import Repository
```

```
repo = Repository("<path_to_tubes-ml>", clone_from="git clone http
```

```
⇒ /usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_depr
For more details, please read https://huggingface.co/docs/huggingfa
warnings.warn(warning_message, FutureWarning)
Cloning https://huggingface.co/rizkiapriliarahman/tubes-ml into loc
WARNING:huggingface_hub.repository:Cloning https://huggingface.co/r
```

Setelah perintah ini dijalankan, repositori dari Hugging Face Hub akan dikloning ke folder lokal yang ditentukan, dan dapat mulai bekerja dengan file di dalamnya, serta mengunggah perubahan atau file baru ke repositori menggunakan `push_to_hub`.

```
[ ] repo.git_pull()  
    repo.git_add()  
    repo.git_push()
```

```
⇒ Everything up-to-date
```

```
WARNING:huggingface_hub.repository:Everything up-to-date
```

```
'https://huggingface.co/rizkiapriliarahman/tubes-ml/commit/1f14f4edd  
669dd76f06956abf7a'
```

Perintah `repo.git_pull()`, `repo.git_add()`, dan `repo.git_push()` digunakan untuk melakukan operasi Git pada repositori lokal yang telah di kloning dari Hugging Face Hub.

```
[ ] repo.git_pull()
```

```
[ ] model.save_pretrained("<path_to_tubes-ml>")  
    tokenizer.save_pretrained("<path_to_tubes-ml>")
```

```
⇒ ('<path_to_tubes-ml>/tokenizer_config.json',  
    '<path_to_tubes-ml>/special_tokens_map.json',  
    '<path_to_tubes-ml>/sentencepiece.bpe.model',  
    '<path_to_tubes-ml>/added_tokens.json',  
    '<path_to_tubes-ml>/tokenizer.json')
```

Perintah `model.save_pretrained("<path_to_tubes-ml>")` dan `tokenizer.save_pretrained("<path_to_tubes-ml>")` digunakan untuk menyimpan model dan tokenizer ke folder lokal. disini perlu mengganti "<path_to_tubes-ml>" dengan path lokal yang diinginkan di system.

```
[ ] repo.git_add()
    repo.git_commit("Add model and tokenizer files")
    repo.git_push()

➡ Upload file model.safetensors: 100% ██████████ 422M/422M [00:09<
    Upload file sentencepiece.bpe.model: 100% ██████████ 792k/792k [00:09<
    To https://huggingface.co/rizkiapriliarahman/tubes-ml
    1f14f4e..8672aa7 main -> main

    WARNING:huggingface_hub.repository:To https://huggingface.co/rizkiap
    1f14f4e..8672aa7 main -> main

    'https://huggingface.co/rizkiapriliarahman/tubes-ml/commit/8672aa7d8
    f1ac88c5199a7ff249'
```

Perintah `repo.git_add()`, `repo.git_commit()`, dan `repo.git_push()` digunakan untuk mengelola perubahan pada repositori Git dan mendorong perubahan ke repositori Hugging Face Hub.

```
[ ] from transformers import AutoModelForMaskedLM, AutoTokenizer

    checkpoint = "camembert-base"

    model = AutoModelForMaskedLM.from_pretrained(checkpoint)
    tokenizer = AutoTokenizer.from_pretrained(checkpoint)

    # Do whatever with the model, train it, fine-tune it...

    model.save_pretrained("<path_to_tubes-ml>")
    tokenizer.save_pretrained("<path_to_tubes-ml>")

➡ Some weights of the model checkpoint at camembert-base were not used
- This IS expected if you are initializing CamembertForMaskedLM from
- This IS NOT expected if you are initializing CamembertForMaskedLM
  ('<path_to_tubes-ml>/tokenizer_config.json',
   '<path_to_tubes-ml>/special_tokens_map.json',
   '<path_to_tubes-ml>/sentencepiece.bpe.model',
   '<path_to_tubes-ml>/added_tokens.json',
   '<path_to_tubes-ml>/tokenizer.json')
```

Perintah yang ini adalah cara untuk memuat model dan tokenizer dari model `camembert-base`, melakukan proses (seperti pelatihan atau fine-tuning) pada model tersebut, dan kemudian menyimpan model dan tokenizer ke folder lokal.